

GUNABHIRAM BILLA

Data & Analytics Engineer | SQL · dbt · Snowflake · Looker · Python · AI-Augmented Workflows

gunabhiram.billa@gmail.com · (940) 287-4848 · linkedin.com/in/gunabhiram · Michigan

PROFESSIONAL SUMMARY

Data and analytics professional with 3+ years of experience integrating, transforming, and analyzing data from multiple enterprise systems using SQL and modern data engineering techniques (dbt, Snowflake, Python). Proven track record developing interactive dashboards, scorecards, and executive reporting in Power BI, Tableau, Looker, and Hex, and working directly with business stakeholders to define reporting requirements, KPIs, and performance measures. Experienced analyzing operational, product, and CRM/case-management data to support executive and cross-functional decision-making, including a 366-test dbt data quality framework with zero production errors and multiple automated reporting and alerting tools built for enterprise operations teams. Rigorously uses AI coding agents (Claude, ChatGPT, Cursor, GitHub Copilot) to accelerate analysis while reviewing and validating all AI-generated outputs. Strong communicator able to translate technical findings into clear, actionable insight for both technical and executive audiences.

TECHNICAL SKILLS

SQL & dbt: Advanced SQL daily (complex joins, subqueries, CTEs, window functions, query optimization, performance tuning); dbt expertise — models, tests, snapshots, macros, packages, exposures; 366-test production dbt data quality framework

Data Warehouses: Snowflake (warehouses, schemas, roles, performance tuning, clustering keys); experience with cloud data warehouses, semantic layers, dimensional modeling, fact/dimension design, slowly changing dimensions (SCD Type 2)

BI & Dashboarding: Power BI (advanced DAX, time intelligence, star schema, RLS, drillthrough), Tableau, Looker, Hex, Amazon QuickSight; metric definition, KPI marts, semantic layer design, governed self-service analytics

AI-Augmented Workflows: Claude, ChatGPT, Cursor, GitHub Copilot for accelerating SQL, Python, and dbt development; rigorous review and refinement of AI-generated code; iterative prompt design for analysis

Python & Automation: Python (pandas, NumPy, automation, scripting); building reusable transformation logic; CI/CD via Git/GitHub with automated testing on pull requests

Cross-functional & Documentation: Business metric definition, data model documentation, source-to-target mappings, data dictionaries; partnering with product, engineering, operations, and business teams; storytelling with data

PROFESSIONAL EXPERIENCE

GM OnStar — Connected Services Systems Engineer (Contractor)

June 2026 – Present

Data & Analytics Engineering, Connected Services — Michigan

- Built a tiered, multi-cell Python + SQL diagnostic tool for fleet telemetry support analysts to identify why vehicles were not reporting telemetry data — a fast-path decision-tree gate evaluation covering consent, enrollment, connectivity, and push history, backed by deep-dive cells for edge cases, replacing manual querying of multiple backend systems and CRM navigation.
- Extended the diagnostic tool with an AI case-note generation layer that produces ready-to-send, partner-facing case notes from collected diagnostic variables and CRM case context, enforcing consistent format and partner identity requirements.
- Designed and shipped an automated hourly Slack alerting system monitoring Siebel CRM activity for Commercial Ops, notifying the team of new cases, customer replies, and case reassignments; diagnosed a multi-hour data lag in the source view that was causing missed alerts and corrected the lookback window accordingly.
- Diagnosed and resolved a duplicate-alert incident caused by a mid-run crash, redesigning the deduplication approach around per-send incremental watermark writes and cross-alert priority logic to prevent recurrence.

- Built a scheduled Python/Spark SQL export tool producing formatted, partner-facing Excel reports on vehicle health status (healthy vs. unhealthy) for containment programs, using window-function deduplication and CASE-based classification to give partner operations teams a single, reliable source of truth per vehicle.
- Built a self-serve, parameter-driven SQL reporting tool giving support analysts a consolidated, real-time view of their open case queue and email activity, surfacing cases needing a response, stale follow-ups, and full ticket inventories without requiring direct navigation of the source CRM.

Yosi Health — Senior Data Analyst (Analytics Engineering Focus)

October 2023 – June 2026

Production Data Models, Semantic Layers, Metrics, Reporting — New York

- Conducted deep-dive analysis and built internal reports, dashboards, and analytics models that inform key product and business decisions across 6+ healthcare product domains (scheduling, intake, payments, reminders, surveys, telehealth) supporting 25M+ records in a fast-paced environment.
- Partnered with product, engineering, operations, and customer success teams to provide data-driven insights, define how key metrics should be calculated and presented, and translate business needs into clear analytical specifications.
- Built and maintained analytics-ready data models using dbt and Snowflake — including reusable semantic layers, KPI marts, and fact/dimension models that centralize business logic and ensure consistent, trusted reporting across the organization.
- Worked directly with business stakeholders across departments to define reporting requirements, KPIs, and performance measures — creating data dictionaries, metric definitions, and source-to-target mappings that empower self-service analytics while maintaining governance and consistency.
- Built a dbt-based data quality framework with 366 passing tests and zero production errors covering source-to-target reconciliation, anomaly detection, business rule validation, and grain integrity — improving reliability and trust in downstream reporting.
- Used AI coding agents (Claude, ChatGPT, Cursor, GitHub Copilot) daily to accelerate SQL and Python development, then rigorously reviewed, refined, and iterated on AI-generated outputs to ensure correctness, performance, and alignment with business logic.
- Developed Python automation for recurring reporting, validation, and reconciliation workflows, improving operational efficiency by 80% and enabling analysts to focus on higher-value investigation work.
- Investigated data anomalies and metric discrepancies through structured root cause analysis — breaking workflows into stages, identifying exception categories, and presenting findings clearly to technical and non-technical stakeholders.
- Built interactive dashboards, scorecards, and executive reports in Power BI (advanced DAX, RLS, drillthrough), Tableau, Looker, and Hex for executive, operational, and customer-facing audiences; designed KPI marts and semantic layers supporting executive decision-making across operational and product performance data.

EDUCATION

University of North Texas, Denton, TX

GPA: 3.6 / 4.0

Master of Science (M.S.) in Business Analytics — Data Analytics & Data Science